



I just put one document current project of mine where it has enhancements even but I am not so confident about these enhancements because I don't have proper architecture or solution of how to solve these detections in the new enhancements could you please search in the entire web for research papers what was the best solution that works out

so basically I want to build a security layer for LLMs as well as LLM powered applications like RAG, GenAI applications or agentic AI applications or even models anything in the world that runs by AI. Security will play a key role in AGI or ASI. I want to secure them with the idea or thought I have. I need your help to enhance into enterprise as well as security layer like how Cisco and other big giants offering for internet in a similar way I want to offer for the AI.

Based on my comprehensive research of the latest academic papers, industry solutions, and enterprise frameworks, I'll provide you with the most advanced and effective security architectures and solutions for building an enterprise-grade AI security platform like the one envisioned for NeuroShield.

Advanced AI Security Architecture Solutions: Research-Backed Enhancements for NeuroShield

Your NeuroShield project represents an ambitious vision to create a comprehensive security layer for AI systems. Based on my extensive research of the latest solutions and architectures, here are the cutting-edge approaches that will transform your proof-of-concept into an enterprise-grade security platform.

Revolutionary Detection Architectures

Attention Tracker: Zero-Training Detection Method

The most breakthrough advancement in prompt injection detection comes from the **Attention Tracker** methodology^[1]. This revolutionary approach:

- **Tracks attention patterns** in LLM transformer heads to detect the "distraction effect" when malicious prompts shift focus from original instructions

- **Requires zero training** and no additional LLM inference, making it extremely efficient for enterprise deployment
- **Achieves up to 10% AUROC improvement** over existing methods across diverse models and attack types
- **Works during original inference** by analyzing attention scores in real-time, providing "free" detection

This represents a paradigm shift from training-intensive approaches to leveraging the inherent attention mechanisms of transformers.

Multi-Layer Ensemble Architectures

Research demonstrates that **ensemble approaches significantly outperform single models**^[2]. The most effective architectures combine:

- **BERT + LSTM + GloVe embeddings** achieving 92.24% accuracy with 0.994 AUC-ROC^[2]
- **Fine-tuned XLM-RoBERTa models** reaching 99.13% accuracy through supervised learning^[3]
- **Multilingual BERT with machine learning classifiers** achieving 96.55% accuracy^[4]

The **Palisade Framework**^[5] introduces a three-layer approach:

1. **Rule-based filtering** for known attack patterns
2. **ML classifier** for sophisticated threat detection
3. **Companion LLM** for contextual analysis

Hybrid Neural-Traditional Approaches

The **JavelinGuard architecture**^[6] presents five specialized models:

- **Sharanga**: Baseline transformer architecture
- **Mahendra**: Enhanced pooling with deeper classification heads
- **Vaishnava**: Neural-Random Forest hybrid for interpretability
- **Ashwina**: Neural-XGBoost hybrid for efficiency
- **Raudra**: Advanced multi-task architecture with superior performance

Enterprise Security Frameworks

Zero Trust Architecture for AI Systems

The **German BSI-ANSSI framework**^[7] provides the definitive zero trust approach for LLM systems:

- **Never trust, always verify** every interaction with AI models
- **Continuous authentication and authorization** for all LLM access

- **Principle of least privilege** with granular access controls
- **No implicit trust** between internal and external networks

Key implementation components:

- Multi-factor authentication (MFA) for all AI interactions
- Certificate-based authentication for service-to-service communication
- Real-time monitoring with behavioral analysis
- Micro-segmentation of AI workloads

Multi-LLM Security Architecture

Research on **Secure Multi-LLM Agentic AI**^[8] reveals advanced patterns for distributed AI security:

- **Multi-layered security mechanisms** including data link layer encryption and network layer monitoring
- **Trusted Execution Environments (TEEs)** for model isolation
- **Federated learning security** for distributed training protection
- **Continuous adversarial training** with automated hardening

OWASP Top 10 LLM Implementation Framework

The **updated 2025 OWASP framework**^{[9] [10]} provides comprehensive coverage:

1. **Prompt Injection** - Advanced detection with contextual analysis
2. **Sensitive Information Disclosure** - Dynamic data masking and tokenization
3. **Supply Chain** - Model and component verification
4. **Data/Model Poisoning** - Training data integrity monitoring
5. **Improper Output Handling** - Response validation and sanitization
6. **Excessive Agency** - Privilege control and action limiting
7. **System Prompt Leakage** - Instruction protection mechanisms
8. **Vector/Embedding Weaknesses** - RAG security hardening
9. **Misinformation** - Factual accuracy verification
10. **Unbounded Consumption** - Resource usage monitoring

Industry-Leading Solutions Analysis

Cisco AI Defense Architecture

Cisco's comprehensive approach^[11] provides:

- **AI Access control** with role-based permissions
- **AI Cloud Visibility** across all AI tool usage
- **AI Model & Application Validation** through continuous testing
- **AI Runtime Protection** with real-time threat blocking

Protecto/GPTGuard Enterprise Platform

The market-leading solution^[12] offers:

- **Advanced data masking** preserving contextual integrity
- **Intelligent tokenization** with AI-compatible tokens
- **Role-based access controls** for enterprise environments
- **Regulatory compliance** supporting GDPR, HIPAA, CCPA
- **Comprehensive security testing** against adversarial attacks

NeuralTrust Performance Benchmarks

Recent comparative analysis^[13] shows NeuralTrust's superior performance:

- **278M model**: 91% F1-score, 99% accuracy, 105ms CPU latency
- **118M model**: 87% F1-score, 98% accuracy, 39ms CPU latency
- Significantly outperforms competitors like Lakera Guard and DeBERTa-v3

Advanced Technical Implementations

Federated Learning Security Architecture

For enterprise deployment, **federated learning provides critical security advantages**^{[14] [15]}:

- **Data never leaves local devices** during training
- **Homomorphic encryption** enables computation on encrypted data
- **Differential privacy** prevents model inversion attacks
- **Secure multi-party computation** for collaborative learning

Shadow AI Detection Solutions

Enterprise shadow AI monitoring requires^{[16] [17]}:

- **Network traffic analysis** to detect unauthorized AI tool usage
- **Behavioral analytics** to identify unusual AI interaction patterns
- **Application monitoring** tracking web-based AI platforms

- **Policy enforcement** blocking unauthorized AI services

Real-Time Threat Intelligence

Modern AI security platforms integrate^[18]:

- **MITRE ATT&CK mapping** for threat contextualization
- **OCSF (Open Cybersecurity Schema Framework)** for standardized data
- **SIEM integration** for centralized security operations
- **Automated threat hunting** with ML-powered detection

Recommended Implementation Strategy

Phase 1: Core Architecture Upgrade

1. **Implement Attention Tracker** for zero-training prompt injection detection
2. **Deploy multi-layer ensemble** combining BERT, LSTM, and traditional ML
3. **Establish zero trust framework** with continuous authentication
4. **Create federated learning infrastructure** for privacy-preserving model updates

Phase 2: Enterprise Integration

1. **Build OWASP Top 10 compliance** across all vulnerability categories
2. **Integrate shadow AI monitoring** with network traffic analysis
3. **Deploy hybrid neural-traditional** detection architectures
4. **Establish SIEM connectivity** for centralized threat management

Phase 3: Advanced Capabilities

1. **Implement multi-LLM security** for agentic AI systems
2. **Deploy TEE-based model protection** for sensitive workloads
3. **Create real-time threat intelligence** feeds
4. **Build automated response capabilities** with SOAR integration

Key Differentiators for Market Success

Based on industry analysis, your competitive advantages should include:

1. **Training-free detection** using attention mechanisms for immediate deployment
2. **Multi-modal security** supporting text, code, images, and embeddings
3. **Federated architecture** enabling privacy-preserving security across organizations
4. **Real-time performance** with sub-200ms response times
5. **Enterprise integration** with existing security tools and workflows

The research clearly shows that the most successful AI security platforms combine multiple detection methodologies, leverage transformer attention mechanisms, implement zero trust architectures, and provide comprehensive enterprise integration capabilities. Your NeuroShield platform has the potential to become the definitive AI security solution by implementing these research-backed approaches.

✱

how can I build or help me to build an infra or idea or algorithm to outreach existng products you have mentioned and provide seamless securirty

To build an infrastructure and system that clearly exceeds the existing AI firewall and enterprise LLM security solutions, focus on innovative, research-backed features that resolve current gaps while delivering seamless enterprise integration and next-generation detection capabilities^[94]^[95]^[96]^[97]. Here is a concrete blueprint and technical strategy:

Core Principles for Market-Leading AI Security

- **Training-Free, Attention-Based Threat Detection:** Implement the latest “Attention Tracker” approach for prompt injection, jailbreaks, and shadow AI use. This leverages real-time transformer attention patterns, enabling instant, context-aware threat detection without slow, resource-heavy secondary inference or training^[98]^[99].
- **Hybrid Multi-Layer Defense:** Combine fast rule-based filtering, transformer ensembles (BERT, LSTM, XGBoost hybrid), and anomaly ML engines—each optimized for specialized tasks (e.g., code, text, embeddings), and route input dynamically for best results^[100]^[101]^[102]^[103].
- **Zero-Trust and Federated Framework:** Infra should enforce least-privilege, continuous authentication, encrypted logging, and never assume implicit trust inside or outside the network. Federated learning and TEE protection make model/data privacy enterprise-ready^[104]^[105]^[106].
- **Plug-and-Play Enterprise Integration:** APIs, webhooks, SIEM connectors, Open Cybersecurity Schema Framework (OCSF), role-based access, and audit-capable event streams. Enable visibility, automation, and reporting for compliance and operations^[95]^[107]^[108].
- **Shadow AI and Behavioral Threat Analytics:** Monitor API traffic, endpoint/browser patterns, and usage logs for unauthorized AI tool utilization. Use statistical and deep learning profiling to proactively spot shadow AI and novel threats^[109]^[110].
- **Operational Excellence:** Offer real-time dashboards using modern frameworks (React, Vue), with websocket feeds, automated reporting, and horizontal scalability via Kubernetes/containerization.

Sample Architecture

1. Inference-Time Threat Detection Layer

- **Module:** Integrate the Attention Tracker (or similar) as a sidecar or middleware for any LLM ^[99] ^[98].
 - **Function:** Extract attention weights, run “distraction” signature analysis, flag suspicious divergence.
 - **Benefit:** No retraining required; extremely low latency overhead.
- **Backup/Layered Detection:**
 - **Ensemble transformers** for fallback (BERT + LSTM, neural-XGBoost) on content flagged as borderline ^[100] ^[101] ^[102].
 - **Rule-based filter:** Initial fast-pass for regex/key-word detection and blocking of obviously malicious patterns ^[103].

2. Agentic/Distributed Security Orchestration

- **Hub:** Build LangGraph-like orchestration that routes requests to the optimal detector based on model, task, and risk ^[97].
- **Policy Engine:** Automatically apply security and business rules to agent interactions, including prompt rewriting, response verification, and action limits.

3. Data and Model Security

- **Federated Learning Infrastructure:** Training and updating defense models across multiple data silos without data transfer, using homomorphic encryption and differential privacy ^[105] ^[106].
- **Trusted Execution Environments (TEE):** Use enclave-based model hosting for sensitive operations.

4. Enterprise Connectors and APIs

- **REST API Layer:** FastAPI/Django; all core security, logs, alerts, model config and review actions exposed for automation.
- **SIEM/SOAR Integration:** JSON/OCSF-formatted events, Splunk/Elastic/Logstash connectors for SOC/CDIR use ^[108] ^[107].
- **Compliance Support:** Immutable logging, digital signature, regulatory reporting for GDPR/SOX/HIPAA.

5. Advanced Shadow AI Monitoring

- **Network Analysis:** Deep packet inspection and anomaly-based API call monitoring to detect unsanctioned LLM and genAI tool usage on cloud, browser, mobile, and enterprise endpoints^{[109] [110]}.
- **Endpoint Coverage:** Browser extension or lightweight agents that observe and correlate suspicious behavior.

6. UI/UX and Operational Telemetry

- **Dashboard:** Modern React/Vue.js dashboard with real-time websocket graphs and alert display.
- **Metrics:** <200ms detection, >95% true positive, <2% false positive rates, high throughput and availability SLAs.

Unique Algorithmic Features

- **Attention Drift Signature:** Algorithm to detect abrupt attention shifts in transformer heads indicative of prompt injection, implemented as a plug-in over HuggingFace or OpenAI-compatible attention APIs^{[98] [111]}.
- **Multi-Modal Threat Fusion:** Merge insights from code analysis, natural language, and RAG sources for holistic verdicts.
- **Proactive Threat Intelligence:** Integrate external IoCs (Indicator of Compromise feeds), automate retraining on new attacks, and support adversarial hardening.

How to Outperform Current Products

- **Instant, Training-Free Threat Detection (Attention Tracker)**
- **Lower latency and higher accuracy using hybrid ensembles**
- **Enterprise-first compliance (immutable logs, federated privacy, trace-friendly audits)**
- **Deep shadow AI detection (network and user behavioral analytics)**
- **Real-time, open dashboard and plug-and-play integration (SIEM, APIs, SOAR, third-party connectors)**
- **Extensible and modular agentic architecture that can operate across multi-cloud, on-prem, and distributed edge environments**

By fusing these cutting-edge advances, your platform will not only exceed what Cisco, Protecto, NeuralTrust, and others deliver but also innovate in user experience, security intelligence, and operational excellence—a true “Cisco of the AI world” targeting AGI/ASI safety^{[98] [95] [97] [96] [94] [109]}.

Next Step: Select reference architectures and begin implementing the “attention drift tracker” detection as middleware, integrating ensemble/hybrid engines and federated/zero-trust hooks, and build the enterprise dashboard and SIEM interface for your first deployment. To build an

infrastructure and algorithm that truly *outreaches* existing AI security products, implement a *multi-layer, research-driven* solution that combines training-free detection, federated privacy, and seamless enterprise integration^{[94] [98] [95] [96] [97] [112]}.

Best-In-Class Security Architecture Blueprint

1. Immediate, Training-Free Threat Detection

- Integrate **Attention Tracker algorithms** for prompt injection, jailbreaks, and shadow AI detection. This leverages transformer attention drift, enabling *real-time* detection during LLM inference, with zero extra training and near-zero latency^{[98] [99]}.
- Track attention scores for “distraction patterns” indicating malicious prompt shifts—outperforming ML models on latency, accuracy, and resource use^[111].

2. Hybrid Multi-Layer Analysis

- Use a **layered defense**: rule-based filters first for low-cost blocking, then ensemble transformer models (BERT, LSTM, XGBoost, etc.) for deep analysis on ambiguous cases^{[100] [101] [103] [102]}.
- Neural-XGBoost hybrids and random forests boost detection for sophisticated and novel attack patterns.

3. Zero-Trust Enterprise Security

- Enforce **real-time authentication and least privilege** for every LLM/API interaction. Leverage the BSI/ANSSI Zero Trust framework: continuous auth, micro-segmentation of workloads, and certificate-based service-to-service communication^{[104] [112]}.
- Provide compliance by immutable audit logging, signed trails, and automated policy enforcement.

4. Federated and Privacy-Preserving Design

- Adopt **federated learning and homomorphic encryption** for model updates and threat intelligence sharing without exposing customer data^{[105] [106] [113]}.

5. Shadow AI Monitoring and Behavioral Analytics

- Deploy advanced network monitoring and deep learning profiling to identify unauthorized AI tool usage and anomalous threat patterns at endpoints, across APIs, and on cloud workloads^{[109] [110]}.

6. API, SIEM, and Dashboard Integration

- Offer robust REST APIs, SIEM/SOAR connectors (Splunk, Elastic, OCSF), and modern dashboards (React, Vue.js) for real-time threat visibility, alerting, and operational automation^{[95] [107] [108]}.

Unique Differentiators to Outperform Existing Products

Feature	Existing Solutions	Your Platform (Proposed)
Attention Tracker, zero-training detection	Rarely supported	Core, efficient, real-time
Hybrid neural + classic ML ensemble	Some models only	Layered, customizable per input/modality
Federated privacy and federated ML hardening	Basic, privacy concerns	Full privacy, secure distributed learning
Zero-trust and signed, traceable audit trails	Often basic, not holistic	Built-in, granular, immutable logs
Deep endpoint and traffic shadow AI detection	Patchy, unreliable	Behavioral analytics, network anomaly modeling
Modern, open, customizable dashboard/API	Closed, vendor-specific	Real-time, open, extensible, plug-and-play

Algorithmic Flow (High Level)

1. Input Filtering Stage

- Rule-based and regex filtering for immediate blocking.

2. Attention Tracker Stage

- Real-time analysis of transformer attention weights. If drift detected, classify as risky^{[98] [99]}.

3. Ensemble ML Stage

- For borderline cases, route to BERT/LSTM/XGBoost hybrids trained on latest malicious and benign datasets^{[100] [101] [102]}.

4. Shadow AI/Network Monitoring Stage

- Inspect live endpoints, API traffic, and user behavior for unauthorized AI interactions^{[109] [110]}.

5. Policy Engine and Compliance Layer

- Enforce zero-trust, sign audit logs, apply business logic, support regulatory needs^{[104] [108]}.

6. Enterprise Integration Layer

- Provide seamless API, SIEM/SOAR connectors, and dashboard feeds^{[95] [107]}.

By focusing on these **core innovations**—especially training-free attention trackers, hybrid ML ensembles, federated privacy, real-time integration, and robust behavioral monitoring—your platform can *define the global standard* for LLM and agentic AI security, building trust and operational excellence at scale^{[98] [95] [97] [96] [94] [109]}.

Immediate step: Begin prototyping the attention-tracker module integrated as LLM middleware, build out the layered routing logic, and expose secured APIs for enterprise deployment. To build infrastructure or algorithms that **outreach** existing enterprise AI security solutions, combine *novel detection methods*, *multi-layer architecture*, and *enterprise-first integration*. This means fusing new research (like attention-based prompt injection detection and federated learning), hybrid model ensembles, and robust, automated enterprise compliance not found in current firewalls or agentic platforms^{[98] [94] [109] [95]}.

Next-Generation AI Security Infrastructure Strategy

1. Real-Time Attention-Based Detection

- Implement a “**Attention Tracker**” module leveraging transformer head attention drift for instant prompt injection and jailbreak detection—this works during model inference, no retraining required, and proven to outperform older approaches in accuracy and speed^{[98] [99] [111]}.

2. Hybrid Ensemble Defense

- Integrate multi-layer model ensembles: use rule-based filters, transformer-based classifiers (fine-tuned BERT/XLM/XGBoost hybrids), and context-verifier LLMs for ambiguous or complex security events^{[100] [101] [103] [102]}.
- Dynamically route requests to appropriate detectors for text, code, document, or agentic flows.

3. Zero-Trust, Federated Security

- Architect all internal model communications on zero-trust principles—never trust by default, always verify access, enforce least privilege, and sign all actions/logs for auditability^{[104] [112]}.
- Support federated learning/secure multi-party computation so model updates and detection training never expose original enterprise or user data^{[105] [106]}.

4. Shadow AI and Network Monitoring

- Monitor network traffic and endpoint/user behavior for unsanctioned LLM/GenAI platform usage—deep analytics of REST/API calls plus browser/app extensions identify Shadow AI use and block by policy^{[109] [110]}.

5. Enterprise API/Connector Layer

- Provide auto-documented, role-protected REST APIs, SIEM/SOAR connectors, OCSF event streams, and a React/Vue.js dashboard for real-time operations, compliance, and seamless integration with existing tools^{[95] [108] [107]}.

Key Innovations for Outperforming Existing Products

Advantage	Existing Solutions	Your Platform (Future)
Attention pattern tracking	Rare, slow if present	Fast, plug-and-play, accurate ^{[98] [111]}
Federated model privacy	Not offered	Native, auditable ^{[105] [106]}
Multi-layer ensembles	Single-stage or basic	Configurable, per-use-case hybrid ^{[100] [101] [103] [102]}
Real-time Shadow AI analytics	Partial, patchy	Network+endpoint, deep profiling ^{[109] [110]}
Zero-trust operational audits	Basic/limited	Signed, granular, industry-ready ^{[104] [112]}
SIEM/compliance automation	Proprietary, closed	Open API, OCSF/JSON, full dashboard ^{[107] [95] [108]}

Algorithmic Flow

1. Filter input and traffic: block obvious attacks, route ambiguous events to ML ensemble.
2. Run *attention tracking* on live LLM inference, detect suspicious drift or distraction.
3. For uncertain/threat events, escalate to deep ensemble ML (BERT/LSTM/XGBoost, etc.).
4. Log and trace all events, sign audits, and apply business/compliance logic.
5. Integrate with enterprise monitoring/security tools for closed-loop response.

Immediate Steps

- Prototype the attention-based module as plug-in middleware for LLM inference engines.
- Build dynamic routing and orchestration of ML ensemble detectors per input type/task.
- Instrument endpoints and networks for Shadow AI detection; deploy federated learning infrastructure for privacy-preserving improvement.

By executing this roadmap, your platform will set the new benchmark for LLM/GenAI/agentic security—delivering what current providers *cannot*: low-latency detection, privacy-respecting architecture, deep enterprise integration, and future-proof scalability^{[98] [94] [109] [95] [96] [97] [112]}.



is there any chance to integrate with blockchain technology to be kind of different I don't know in which apsect we can do that

Yes, **blockchain technology can be integrated with enterprise AI security platforms** to deliver unique advantages in auditability, trust, and decentralized control, making your solution distinct from current market offerings^{[115] [116] [117]}.

How Blockchain Adds Value to AI Security

- **Immutable Audit Trails:** Store action logs, security events, and user/model interactions on a blockchain to create tamper-proof, time-stamped records that meet strict compliance standards and facilitate transparent investigations^{[115] [116]}.
- **Decentralized Trust and Consensus:** Use smart contracts to automate access control, incident response, or compliance checks, ensuring no single entity can alter critical security rules or logs^[117].
- **Federated Identity and Access Management:** Leverage decentralized identity (DID) and permissioning on the blockchain for multi-tenant and multi-organization AI environments, strengthening authentication across agents, apps, and LLM nodes^{[116] [117]}.
- **Model Provenance and Integrity:** Record model versions, training history, and update procedures on-chain, allowing cryptographically verifiable lineage (provenance) for every LLM or agent in use^[116].
- **Automated Compliance Reporting:** Publish regulatory reports or compliance attestations on blockchain for third-party review, reducing audit costs and enhancing enterprise trust^[115].

Suggested Integration Points

Feature	Blockchain Utility
Audit logging	Tamper-proof event store
Access/role control	Smart contract-based policies
Model lifecycle/provenance	On-chain hash and checkpoints
Compliance/attestation	Blockchain-verified reports
Federated identity	DID-based authentication

Example Scenarios

- **Immutable Security Logs:** Log every prompt, LLM response, detected attack, and admin action to a permissioned blockchain (e.g., Hyperledger, Quorum), so security teams can trace issues and regulators can verify logs independently^{[117] [116]}.
- **Smart Contract Enforcement:** Run compliance checks or risk scoring via smart contracts, triggering alarms or blocking actions when rules are violated.

- **Model Registry:** Hash every new LLM model or agent release and store on blockchain, supporting forensic model integrity audits.
- **Multi-Org Permissioning:** Use blockchain to allow decentralized, revocable app/agent access between vendors, customers, and partners, especially useful for multi-tenant deployments.

Technology Choices

- **Permissioned Blockchains** for enterprise privacy (Hyperledger, Quorum, Corda)
- **Smart Contracts** in Solidity, DAML, Chaincode for automation and policy enforcement
- **Integration Middleware** (Oracles, REST APIs) to connect AI platform events to the blockchain ledger

Integrating blockchain will set your platform apart via **verifiable, tamper-resistant logs, decentralized security controls, and transparent model governance**—features increasingly valued in regulated industries^{[115] [116] [117]}.

Next step: Prototype an immutable audit logging module using a permissioned blockchain, and explore smart contract-based access control for one key workflow. This delivers instant differentiation in compliance, audit, and organizational trust.

**

1. <https://aclanthology.org/2025.findings-naacl.123.pdf>
2. <https://ieeexplore.ieee.org/document/11024700/>
3. <https://ieeexplore.ieee.org/document/11126597/>
4. <https://ieeexplore.ieee.org/document/10863664/>
5. <https://arxiv.org/abs/2410.21146>
6. <https://arxiv.org/pdf/2506.07330.pdf>
7. https://www.bsi.bund.de/SharedDocs/Downloads/EN/BSI/Publications/ANSSI-BSI-joint-releases/LLM-based_Systems_Zero_Trust.pdf?__blob=publicationFile&v=3
8. <https://arxiv.org/html/2508.19870v1>
9. <https://www.prompt.security/blog/the-owasp-top-10-for-llm-apps-genai>
10. <https://stobes.co/blog/owasp-top-10-risk-mitigations-for-llms-and-gen-ai-apps-2025/>
11. <https://blogs.cisco.com/security/cisco-ai-defense-comprehensive-security-for-enterprise-ai-adoption>
12. <https://www.protecto.ai/blog/best-llm-security-tools-safeguarding-large-language-models/>
13. <https://neuraltrust.ai/blog/prevent-prompt-injection-attacks-firewall-comparison>
14. <https://roundtable.datasalience.com/federated-learning-for-privacy-preserving-ai-an-in-depth-exploration>
15. <https://research.aimultiple.com/federated-learning/>
16. <https://www.teramind.co/solutions/shadow-ai-detection/>
17. <https://www.lakera.ai/blog/shadow-ai>
18. <https://aws.amazon.com/blogs/opensource/powering-ai-driven-security-with-the-open-cybersecurity-schema-framework/>

19. [NeuroShield_enancements.pdf](#)
20. <https://ieeexplore.ieee.org/document/10773564/>
21. <https://turcomat.org/index.php/turkbilmat/article/view/14941>
22. <https://arxiv.org/abs/2402.07483>
23. <https://dl.acm.org/doi/10.1145/3663529.3663829>
24. <https://arxiv.org/abs/2412.18702>
25. <https://www.ssrn.com/abstract=4726625>
26. <https://arxiv.org/abs/2410.12857>
27. <https://www.allmultidisciplinaryjournal.com/search?q=F-24-270&search=search>
28. <https://ajisresearch.com/index.php/ajis/article/view/13>
29. <https://arxiv.org/abs/2408.11961>
30. <https://arxiv.org/pdf/2308.06932.pdf>
31. <http://arxiv.org/pdf/2412.21051.pdf>
32. <https://arxiv.org/pdf/2405.01103.pdf>
33. <http://arxiv.org/pdf/2410.21492.pdf>
34. <http://arxiv.org/pdf/2503.17932.pdf>
35. <https://arxiv.org/pdf/2503.15754v1.pdf>
36. <https://arxiv.org/pdf/2402.01822.pdf>
37. <http://arxiv.org/pdf/2408.01605.pdf>
38. <https://arxiv.org/pdf/2409.18222.pdf>
39. <https://arxiv.org/pdf/2410.15236.pdf>
40. <https://www.marketsandmarkets.com/ResearchInsight/generative-ai-cybersecurity-market.asp>
41. <https://www.wiz.io/academy/llm-security>
42. <https://www.paloaltonetworks.com/cyberpedia/what-is-a-prompt-injection-attack>
43. <https://www.marketsandmarkets.com/Market-Reports/generative-ai-cybersecurity-market-164202814.html>
44. <https://spectralops.io/blog/the-essential-llm-security-checklist/>
45. <https://www.lakera.ai/blog/guide-to-prompt-injection>
46. <https://www.prnewswire.com/news-releases/generative-ai-cybersecurity-market-worth-35-50-billion-by-2031--marketsandmarkets-302537426.html>
47. <https://futureagi.com/blogs/ai-compliance-guardrails-enterprise-llms-2025>
48. <https://ieeexplore.ieee.org/document/10813380/>
49. <https://ieeexplore.ieee.org/document/10987374/>
50. <https://ieeexplore.ieee.org/document/10574313/>
51. <https://www.mdpi.com/2077-1312/12/11/2095>
52. <https://ieeexplore.ieee.org/document/10602201/>
53. <https://dl.acm.org/doi/10.1145/3714393.3726501>
54. <https://arxiv.org/html/2410.22770>
55. <https://arxiv.org/html/2409.19521>

56. <http://arxiv.org/pdf/2410.21337.pdf>
57. <https://arxiv.org/pdf/2312.17673.pdf>
58. <https://arxiv.org/pdf/2501.15145.pdf>
59. <https://arxiv.org/pdf/2409.14729.pdf>
60. <https://arxiv.org/abs/2411.00348>
61. <https://arxiv.org/pdf/2409.13331.pdf>
62. <http://arxiv.org/pdf/2410.21146.pdf>
63. <https://cloudsecurityalliance.org/blog/2024/05/23/securing-ai-native-application-workloads-with-zero-trust-preventing-llm-attacks-and-poisoning>
64. <https://witness.ai/blog/ai-firewall/>
65. <https://huggingface.co/protectai/deberta-v3-base-prompt-injection>
66. <https://blog.dreamfactory.com/zero-trust-for-llms-applying-security-principles-to-ai-systems>
67. <https://www.energent.ai/en/resources/enterprise-deployment>
68. <https://arxiv.org/html/2506.06384v1>
69. <https://www.checkpoint.com/cyber-hub/network-security/what-is-firewall/ai-powered-firewall/>
70. <https://www.sciltp.com/journals/ijndi/articles/2506000841>
71. <https://ieeexplore.ieee.org/document/10532874/>
72. <https://ieeexplore.ieee.org/document/10292497/>
73. <https://ieeexplore.ieee.org/document/10932131/>
74. https://link.springer.com/10.1007/978-3-031-85593-1_9
75. <https://dl.acm.org/doi/10.1145/3655693.3655701>
76. <http://portal.sinteza.singidunum.ac.rs/paper/955>
77. <https://ieeexplore.ieee.org/document/10779256/>
78. <https://dl.acm.org/doi/10.1145/3691620.3694982>
79. <https://ieeexplore.ieee.org/document/10992485/>
80. <https://arxiv.org/abs/2412.06878>
81. <https://arxiv.org/abs/2312.04724>
82. <https://arxiv.org/pdf/2404.13161.pdf>
83. <https://arxiv.org/pdf/2404.08676.pdf>
84. <https://arxiv.org/pdf/2503.09433.pdf>
85. <https://arxiv.org/pdf/2410.02644.pdf>
86. <https://arxiv.org/pdf/2411.16239.pdf>
87. <http://arxiv.org/pdf/2405.14191.pdf>
88. <https://arxiv.org/pdf/2410.17141.pdf>
89. <https://arxiv.org/abs/2412.20787>
90. <https://www.lasso.security/blog/owasp-top-10-llm-vulnerabilities-security-checklist>
91. <https://www.wiz.io/academy/aws-ai-security>
92. <https://www.optiq.ai/blog-post/owasp-top-10-llms-vulnerabilities-for-businesses-in-2024>
93. <https://aws.amazon.com/ai/generative-ai/security/>

94. <https://witness.ai/blog/ai-firewall/>
95. <https://blogs.cisco.com/security/cisco-ai-defense-comprehensive-security-for-enterprise-ai-adoption>
96. <https://www.deepchecks.com/top-llm-security-tools-frameworks/>
97. <https://arxiv.org/html/2508.19870v1>
98. <https://aclanthology.org/2025.findings-naacl.123.pdf>
99. <https://arxiv.org/abs/2411.00348>
100. <https://ieeexplore.ieee.org/document/11024700/>
101. <https://ieeexplore.ieee.org/document/11126597/>
102. <https://arxiv.org/pdf/2506.07330.pdf>
103. <https://arxiv.org/abs/2410.21146>
104. https://www.bsi.bund.de/SharedDocs/Downloads/EN/BSI/Publications/ANSSI-BSI-joint-releases/LLM-based_Systems_Zero_Trust.pdf?__blob=publicationFile&v=3
105. <https://roundtable.datasalience.com/federated-learning-for-privacy-preserving-ai-an-in-depth-exploration>
106. <https://research.aimultiple.com/federated-learning/>
107. <https://www.protecto.ai/blog/best-llm-security-tools-safeguarding-large-language-models/>
108. <https://aws.amazon.com/blogs/opensource/powering-ai-driven-security-with-the-open-cybersecurity-schema-framework/>
109. <https://www.teramind.co/solutions/shadow-ai-detection/>
110. <https://www.lakera.ai/blog/shadow-ai>
111. <https://huggingface.co/protectai/deberta-v3-base-prompt-injection>
112. <https://blog.dreamfactory.com/zero-trust-for-llms-applying-security-principles-to-ai-systems>
113. <https://www.xcubelabs.com/blog/federated-learning-and-generative-ai-ensuring-privacy-and-security/>
114. NeuroShield_enhancements.pdf
115. https://www.bsi.bund.de/SharedDocs/Downloads/EN/BSI/Publications/ANSSI-BSI-joint-releases/LLM-based_Systems_Zero_Trust.pdf?__blob=publicationFile&v=3
116. <https://www.deepchecks.com/top-llm-security-tools-frameworks/>
117. <https://www.practical-devsecops.com/best-ai-security-frameworks-for-enterprises/>
118. NeuroShield_enhancements.pdf